

# Simplificarea execuției kernel-urilor OpenCL și studiu de performanță pe diverse platforme hardware

Student: Vlad STOEAN

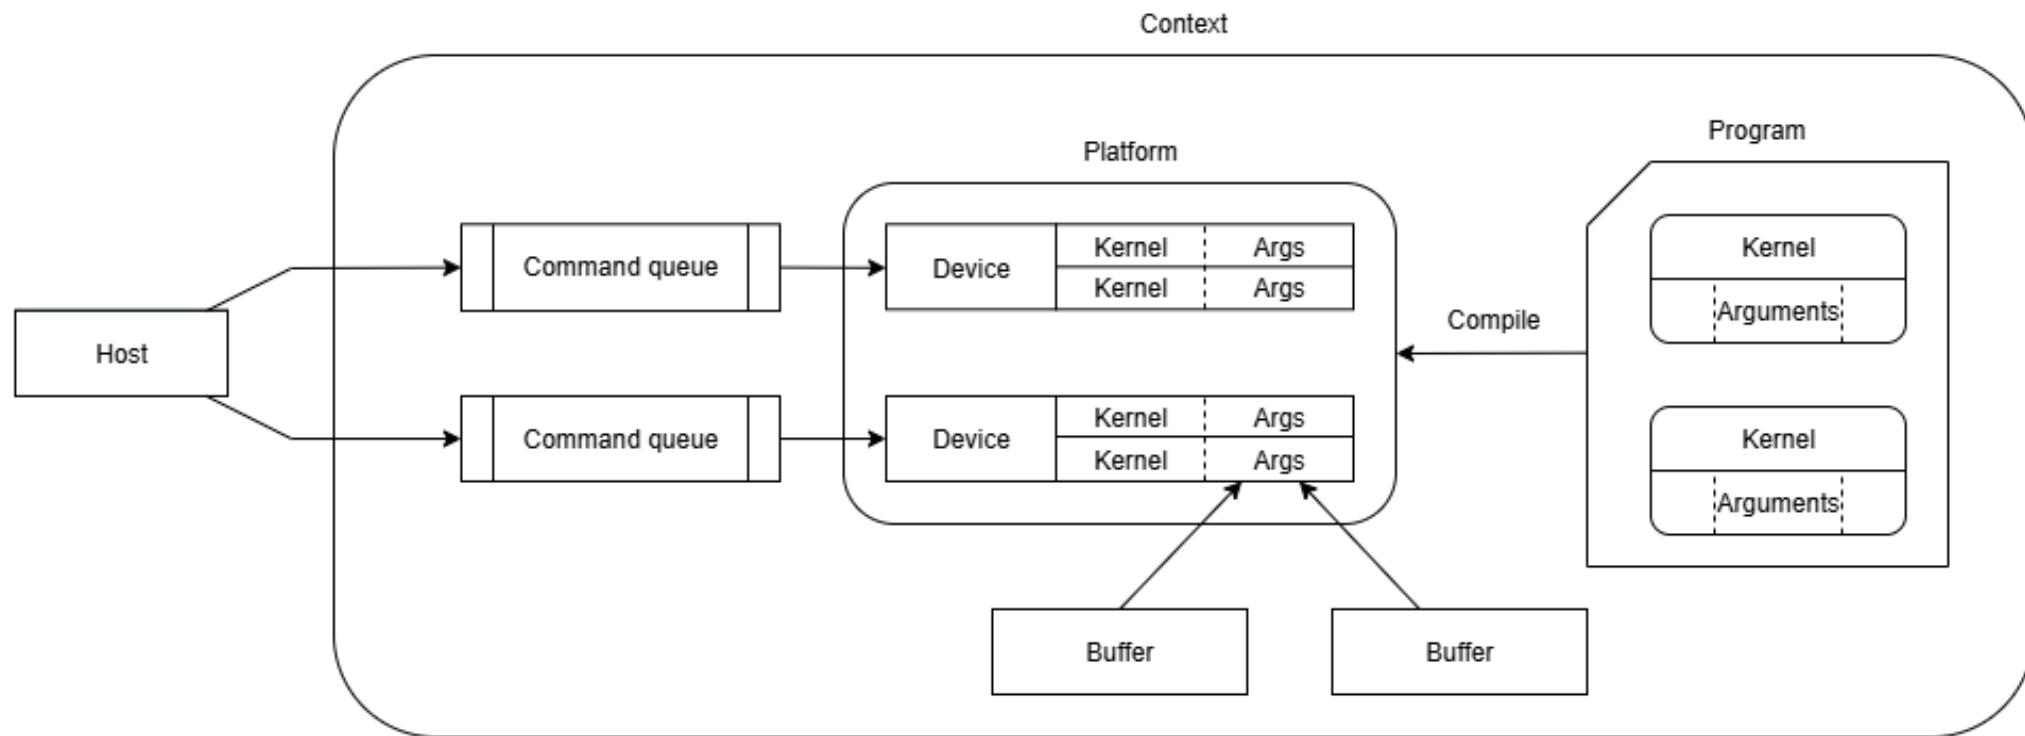
Profesor Îndrumător: conf. dr. ing. Andrei STAN



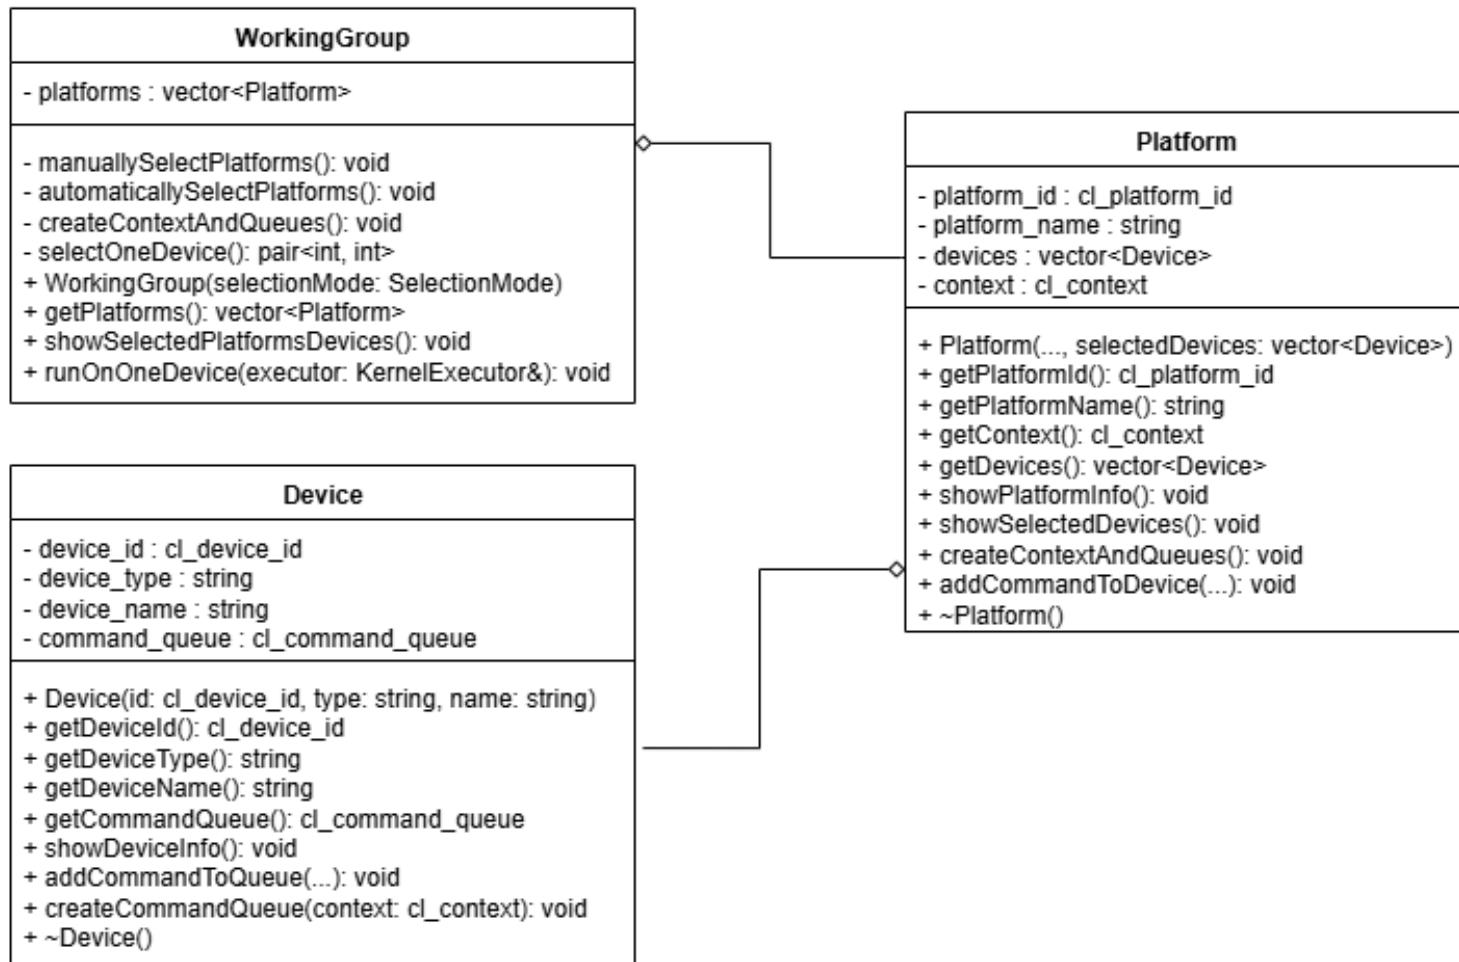
# Contextul și Motivarea Temei

Acceleratoarele embedded sunt utilizate tot mai frecvent în aplicații moderne precum: mașini autonome, supraveghere video intelligentă și drone autonome. Acestea necesită procesare paralelă rapidă pentru analiză de date în timp real.

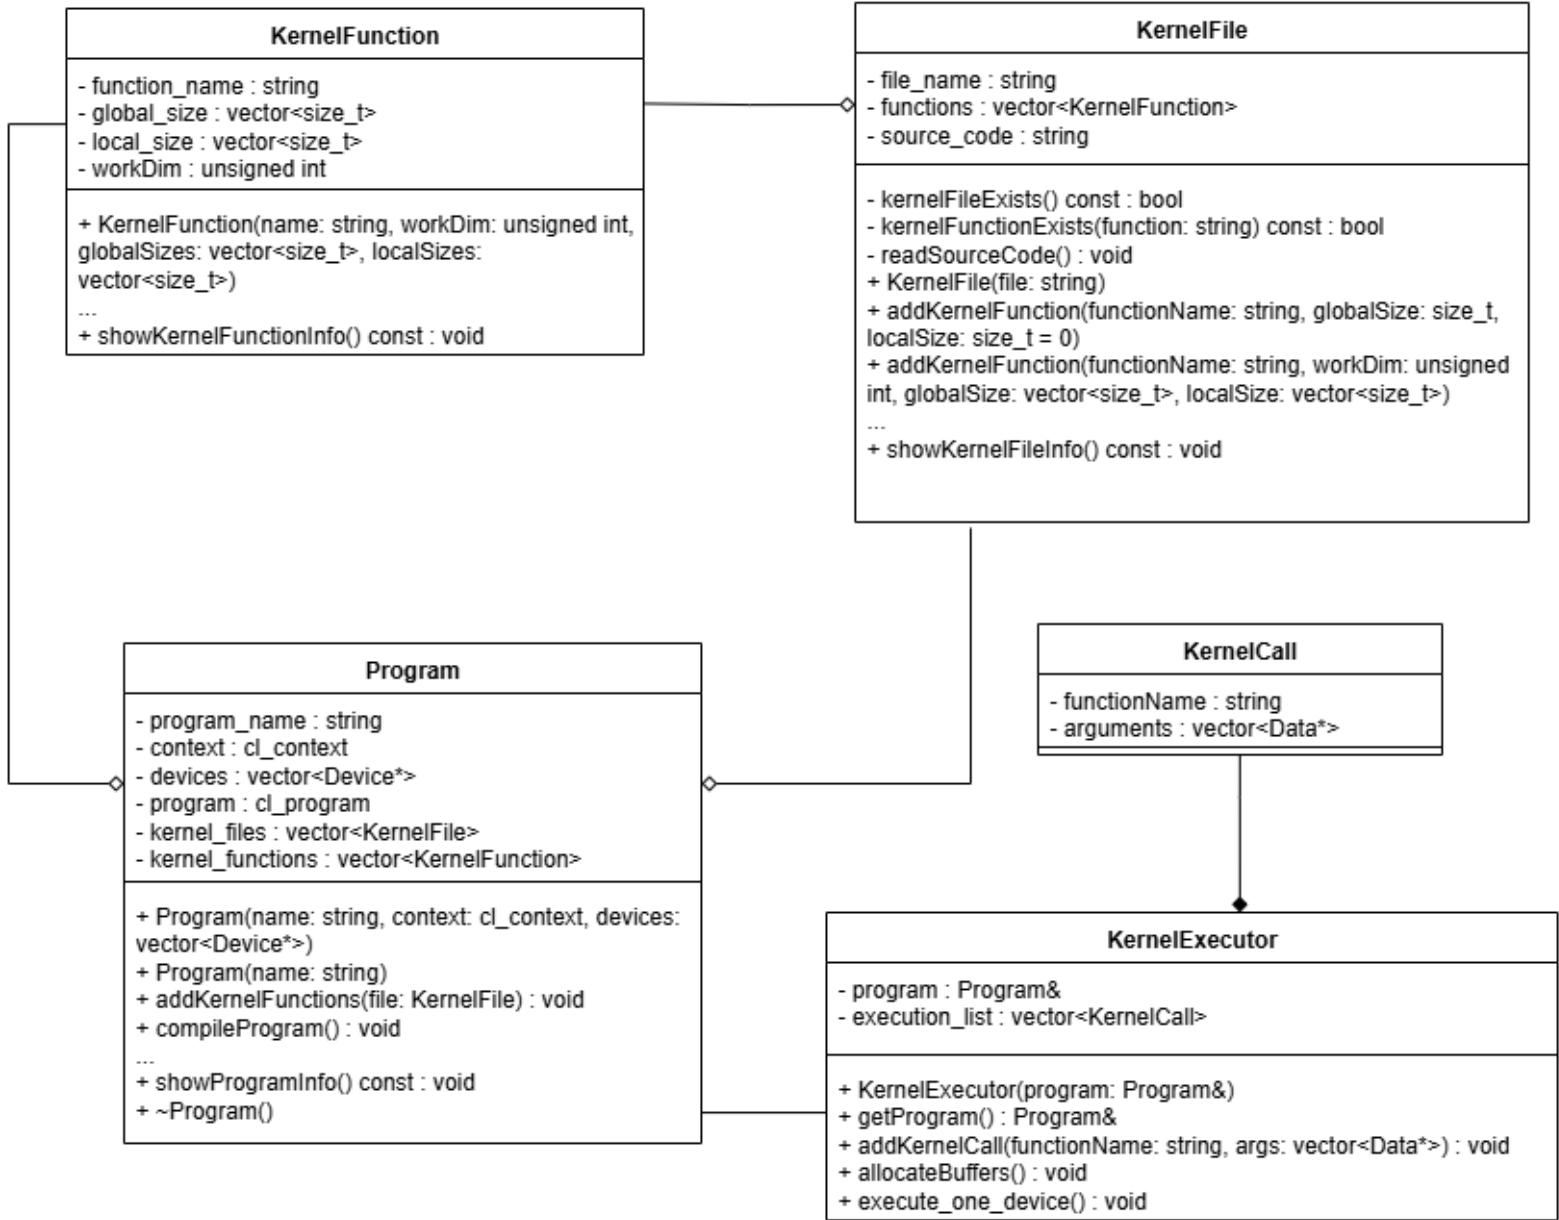
# Modelul de lucru OpenCL



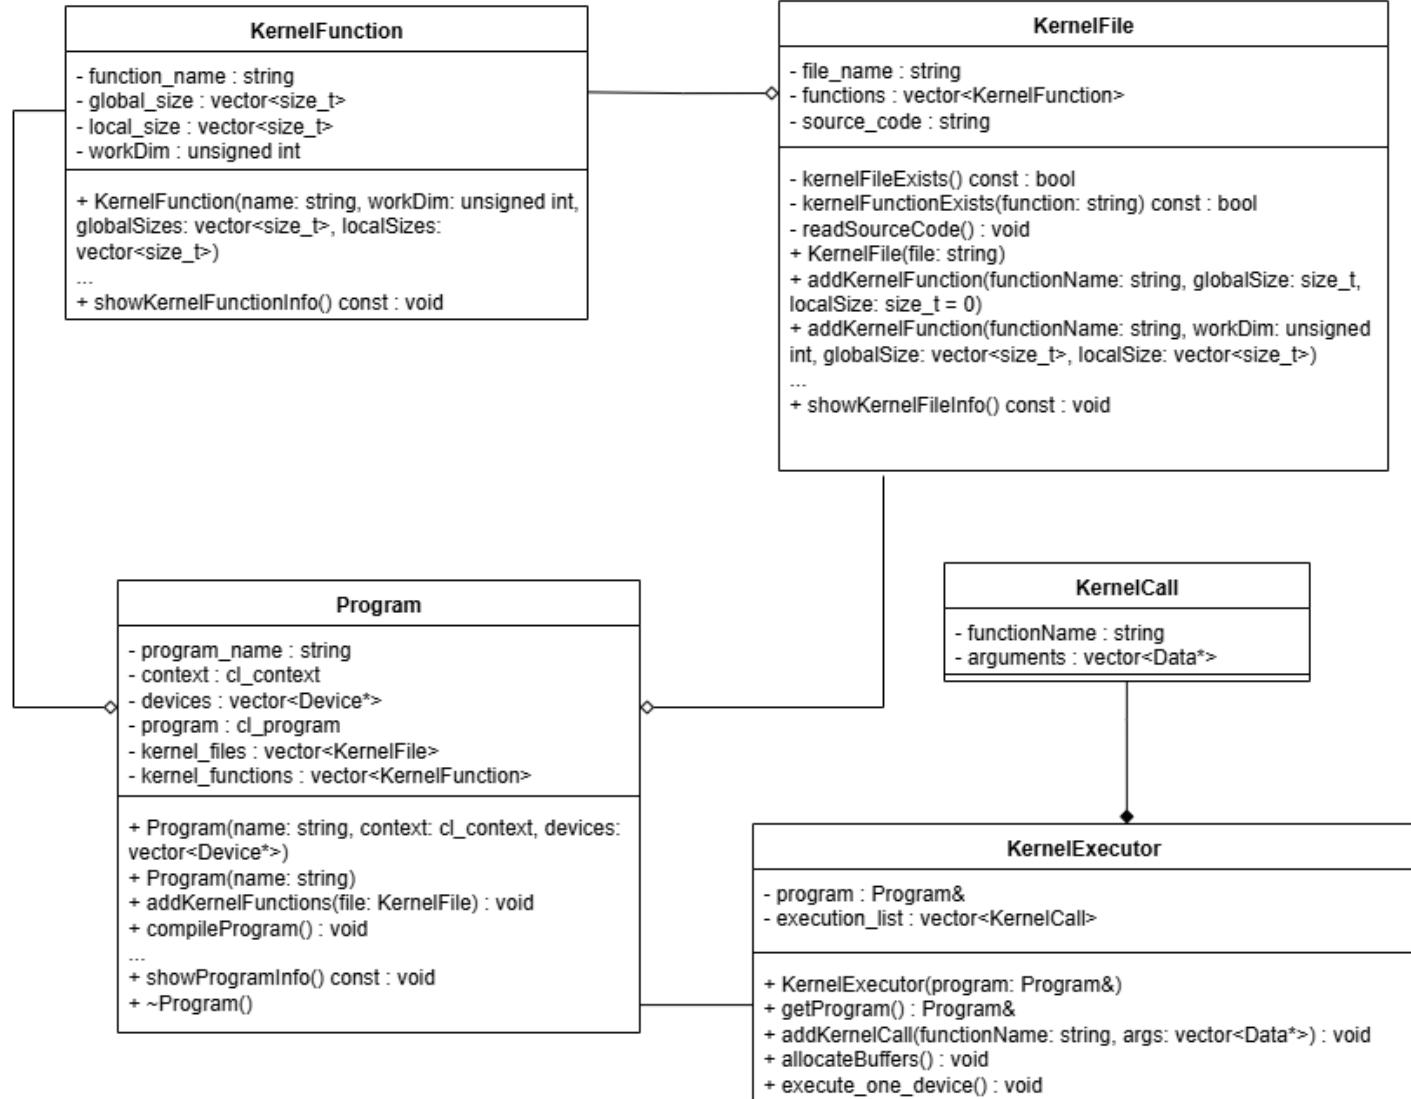
# Clase pentru identificarea și alegerea resurselor hardware



# Clase pentru definirea și execuția kernelurilor OpenCL



# Clase pentru gestionarea argumentelor kernel-urilor



# Clasă pentru scrierea și citirea datelor de test

## Utils

```
+ readVectorFromFile<T>(parentFolder: string, filename: string): vector<T>
+ writeVectorToFile<T>(parentFolder: string, filename: string, data: vector<T>): void
+ readMatrixFromFile<T>(parentFolder: string, filename: string, rows: unsigned int&, cols: unsigned int&): vector<T>
± generateRandomMatrixToFile(folder: string, filename: string, rows: int, cols: int): void
```



LOC folosind framework-ul: 23

```
// Device selection
WorkingGroup workingGroup(SelectionMode::Manual);

// KernelFiles
KernelFile vectorAdd("vector_add.cl");

// Host memory initialization
vector<float> a_vec = Utils::readVectorFromFile<float>("vector_add_f32", "a.txt");
vector<float> b_vec = Utils::readVectorFromFile<float>("vector_add_f32", "b.txt");

// KernelFunctions
vectorAdd.addKernelFunction("vector_add_f32", 4096);

// Program initialization
Program program("VectorOps");
program.addKernelFunctions(vectorAdd);

// OpenCL Data - Buffers and Scalar
Buffer a = Buffer::fromValues<float>(a_vec, Access::ReadOnly);
Buffer b = Buffer::fromValues<float>(b_vec, Access::ReadOnly);
Buffer c = Buffer::empty<float>(4096, Access::ReadWrite);

// KernelExecutor - program flow
KernelExecutor executor(program);
executor.addKernelCall("vector_add_f32", {&a, &b, &c});

// Running
workingGroup.runOnOneDevice(executor);

// Result
vector<float> result = c.readBack<float>(program.getDevices()[0]->getCommandQueue());
Utils::writeVectorToFile<float>("vector_add_f32", "c.txt", result);
```

LOC fără a folosi framework-ul: 134

Implementarea directă în OpenCL  
necesită aproximativ de 6 ori mai  
mult cod decât utilizarea  
framework-ului.

```
// Device selection
cl_uint numPlatforms = 0;
clGetPlatformIDs(0, nullptr, &numPlatforms);
vector<cl_platform_id> platforms(numPlatforms);
clGetPlatformIDs(numPlatforms, platforms.data(), nullptr);

vector<cl_device_id> allDevices;
vector<string> deviceNames;

for (auto platform : platforms) {
    cl_uint numDevices = 0;
    clGetDeviceIDs(platform, CL_DEVICE_TYPE_ALL, 0, nullptr, &numDevices);
    vector<cl_device_id> devices(numDevices);
    clGetDeviceIDs(platform, CL_DEVICE_TYPE_ALL, numDevices, devices.data(), nullptr);

    for (auto device : devices) {
        size_t nameSize = 0;
        clGetDeviceInfo(device, CL_DEVICE_NAME, 0, nullptr, &nameSize);
        vector<char> nameBuffer(nameSize);
        clGetDeviceInfo(device, CL_DEVICE_NAME, nameSize, nameBuffer.data(), nullptr);
        string deviceName(nameBuffer.data());

        deviceNames.push_back(deviceName);
        allDevices.push_back(device);
    }
}

cout << "Choose a device:\n";
for (int i = 0; i < deviceNames.size(); ++i)
    cout << "[" << i << "] " << deviceNames[i] << "\n";

int selection = 0;
cout << "Index: ";
cin >> selection;

cl_device_id selectedDevice = allDevices[selection];
cl_int err;
cl_context context = clCreateContext(nullptr, 1, &selectedDevice, nullptr, nullptr, &err);
if (!context || err != CL_SUCCESS)
    throw runtime_error("Failed to create OpenCL context.");

cl_command_queue queue = clCreateCommandQueue(context, selectedDevice, CL_QUEUE_PROFILING_ENABLE, &err);
if (!queue || err != CL_SUCCESS)
    throw runtime_error("Failed to create command queue.");

// KernelFiles
ifstream sourceFile("kernels/vector_add.cl");
if (!sourceFile.is_open())
    throw runtime_error("Could not open vector_add.cl file");

stringstream sourceBuffer;
sourceBuffer << sourceFile.rdbuf();
string sourceCode = sourceBuffer.str();
const char* sourceStr = sourceCode.c_str();

// Host memory initialization
vector<float> a_vec, b_vec;
{
    ifstream in_a("data/vector_add_f32/a.txt");
    throw runtime_error("Could not open a.txt");
    float val;
    while (in_a >> val) a_vec.push_back(val);
}

{
    ifstream in_b("data/vector_add_f32/b.txt");
    throw runtime_error("Could not open b.txt");
    float val;
    while (in_b >> val) b_vec.push_back(val);
}

// KernelFunctions
const char* kernelName = "vector_add_f32";
size_t globalSize = 4096;
```

```
// Program initialization
cl_program program = clCreateProgramWithSource(context, 1, &sourceStr, nullptr, nullptr);
err = clBuildProgram(program, 1, &selectedDevice, nullptr, nullptr, nullptr);
if (err != CL_SUCCESS) {
    size_t logSize;
    clGetProgramBuildInfo(program, selectedDevice, CL_PROGRAM_BUILD_LOG, 0, nullptr, &logSize);
    vector<char> log(logSize);
    clGetProgramBuildInfo(program, selectedDevice, CL_PROGRAM_BUILD_LOG, logSize, log.data(), nullptr);
    cerr << "Build error:\n" << log.data() << '\n';
    throw runtime_error("OpenCL build failed.");
}

// OpenCL Data - Buffers and Scalar
cl_mem bufA = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(float) * a_vec.size(), nullptr, nullptr);
cl_mem bufB = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(float) * b_vec.size(), nullptr, nullptr);
cl_mem bufC = clCreateBuffer(context, CL_MEM_READ_WRITE, sizeof(float) * 4096, nullptr, nullptr);

cl_event writeEventA, writeEventB;
clEnqueueWriteBuffer(queue, bufA, CL_TRUE, 0, sizeof(float) * a_vec.size(), a_vec.data(), 0, nullptr, &writeEventA);
clEnqueueWriteBuffer(queue, bufB, CL_TRUE, 0, sizeof(float) * b_vec.size(), b_vec.data(), 0, nullptr, &writeEventB);

cl_ulong start, end;
double writeTimeA, writeTimeB;

clGetEventProfilingInfo(writeEventA, CL_PROFILING_COMMAND_START, sizeof(cl_ulong), &start, nullptr);
clGetEventProfilingInfo(writeEventA, CL_PROFILING_COMMAND_END, sizeof(cl_ulong), &end, nullptr);
writeTimeA = (end - start) * 1e-6;

clGetEventProfilingInfo(writeEventB, CL_PROFILING_COMMAND_START, sizeof(cl_ulong), &start, nullptr);
clGetEventProfilingInfo(writeEventB, CL_PROFILING_COMMAND_END, sizeof(cl_ulong), &end, nullptr);
writeTimeB = (end - start) * 1e-6;

cout << "Write time A: " << writeTimeA << " ms\n";
cout << "Write time B: " << writeTimeB << " ms\n";
// KernelExecutor - program flow
cl_kernel kernel = clCreateKernel(program, kernelName, nullptr);

clSetKernelArg(kernel, 0, sizeof(cl_mem), &bufA);
clSetKernelArg(kernel, 1, sizeof(cl_mem), &bufB);
clSetKernelArg(kernel, 2, sizeof(cl_mem), &bufC);

// Running
double kernelTime, readTime;

cl_event execEvent;
clEnqueueNDRangeKernel(queue, kernel, 1, nullptr, &globalSize, nullptr, 0, nullptr, &execEvent);
clFinish(queue);

clGetEventProfilingInfo(execEvent, CL_PROFILING_COMMAND_START, sizeof(cl_ulong), &start, nullptr);
clGetEventProfilingInfo(execEvent, CL_PROFILING_COMMAND_END, sizeof(cl_ulong), &end, nullptr);
kernelTime = (end - start) * 1e-6;

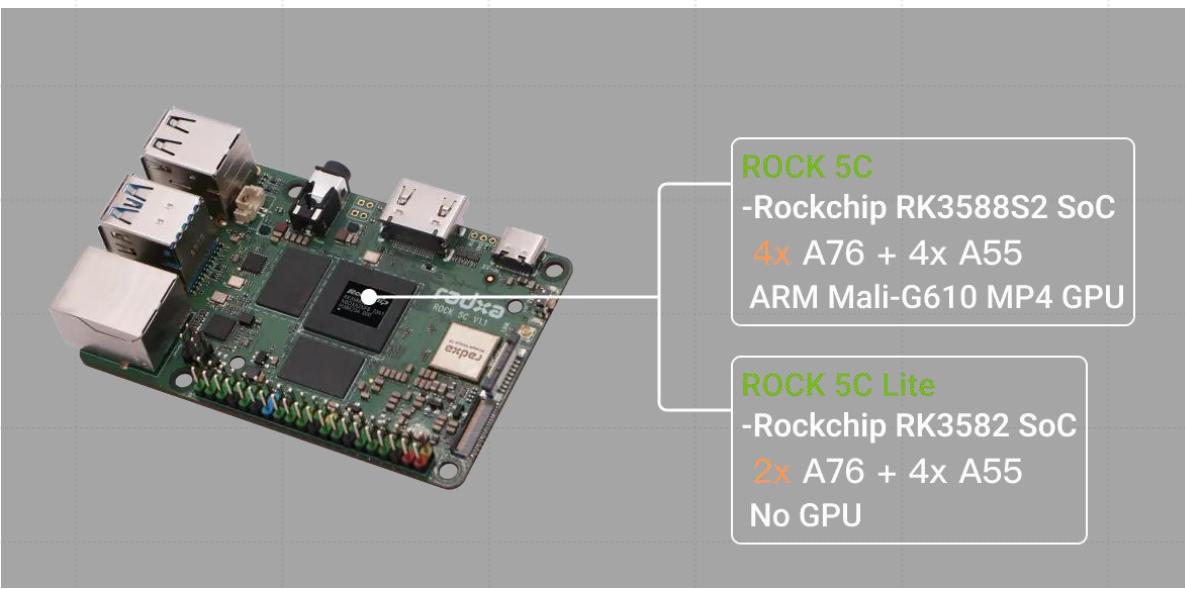
// Result
vector<float> result(4096);
cl_event readEvent;
clEnqueueReadBuffer(queue, bufC, CL_TRUE, 0, sizeof(float) * 4096, result.data(), 0, nullptr, &readEvent);
clFinish(queue);

clGetEventProfilingInfo(readEvent, CL_PROFILING_COMMAND_START, sizeof(cl_ulong), &start, nullptr);
clGetEventProfilingInfo(readEvent, CL_PROFILING_COMMAND_END, sizeof(cl_ulong), &end, nullptr);
readTime = (end - start) * 1e-6;

cout << "Kernel execution time: " << kernelTime << " ms\n";
cout << "Readback time: " << readTime << " ms\n";

ofstream out("data/vector_add_f32/c.txt");
for (float val : result)
    out << val << '\n';
```

# Platformele folosite pentru testarea framework-ului



## ROCK 5C

-Rockchip RK3588S2 SoC  
4x A76 + 4x A55  
ARM Mali-G610 MP4 GPU

## ROCK 5C Lite

-Rockchip RK3582 SoC  
2x A76 + 4x A55  
No GPU





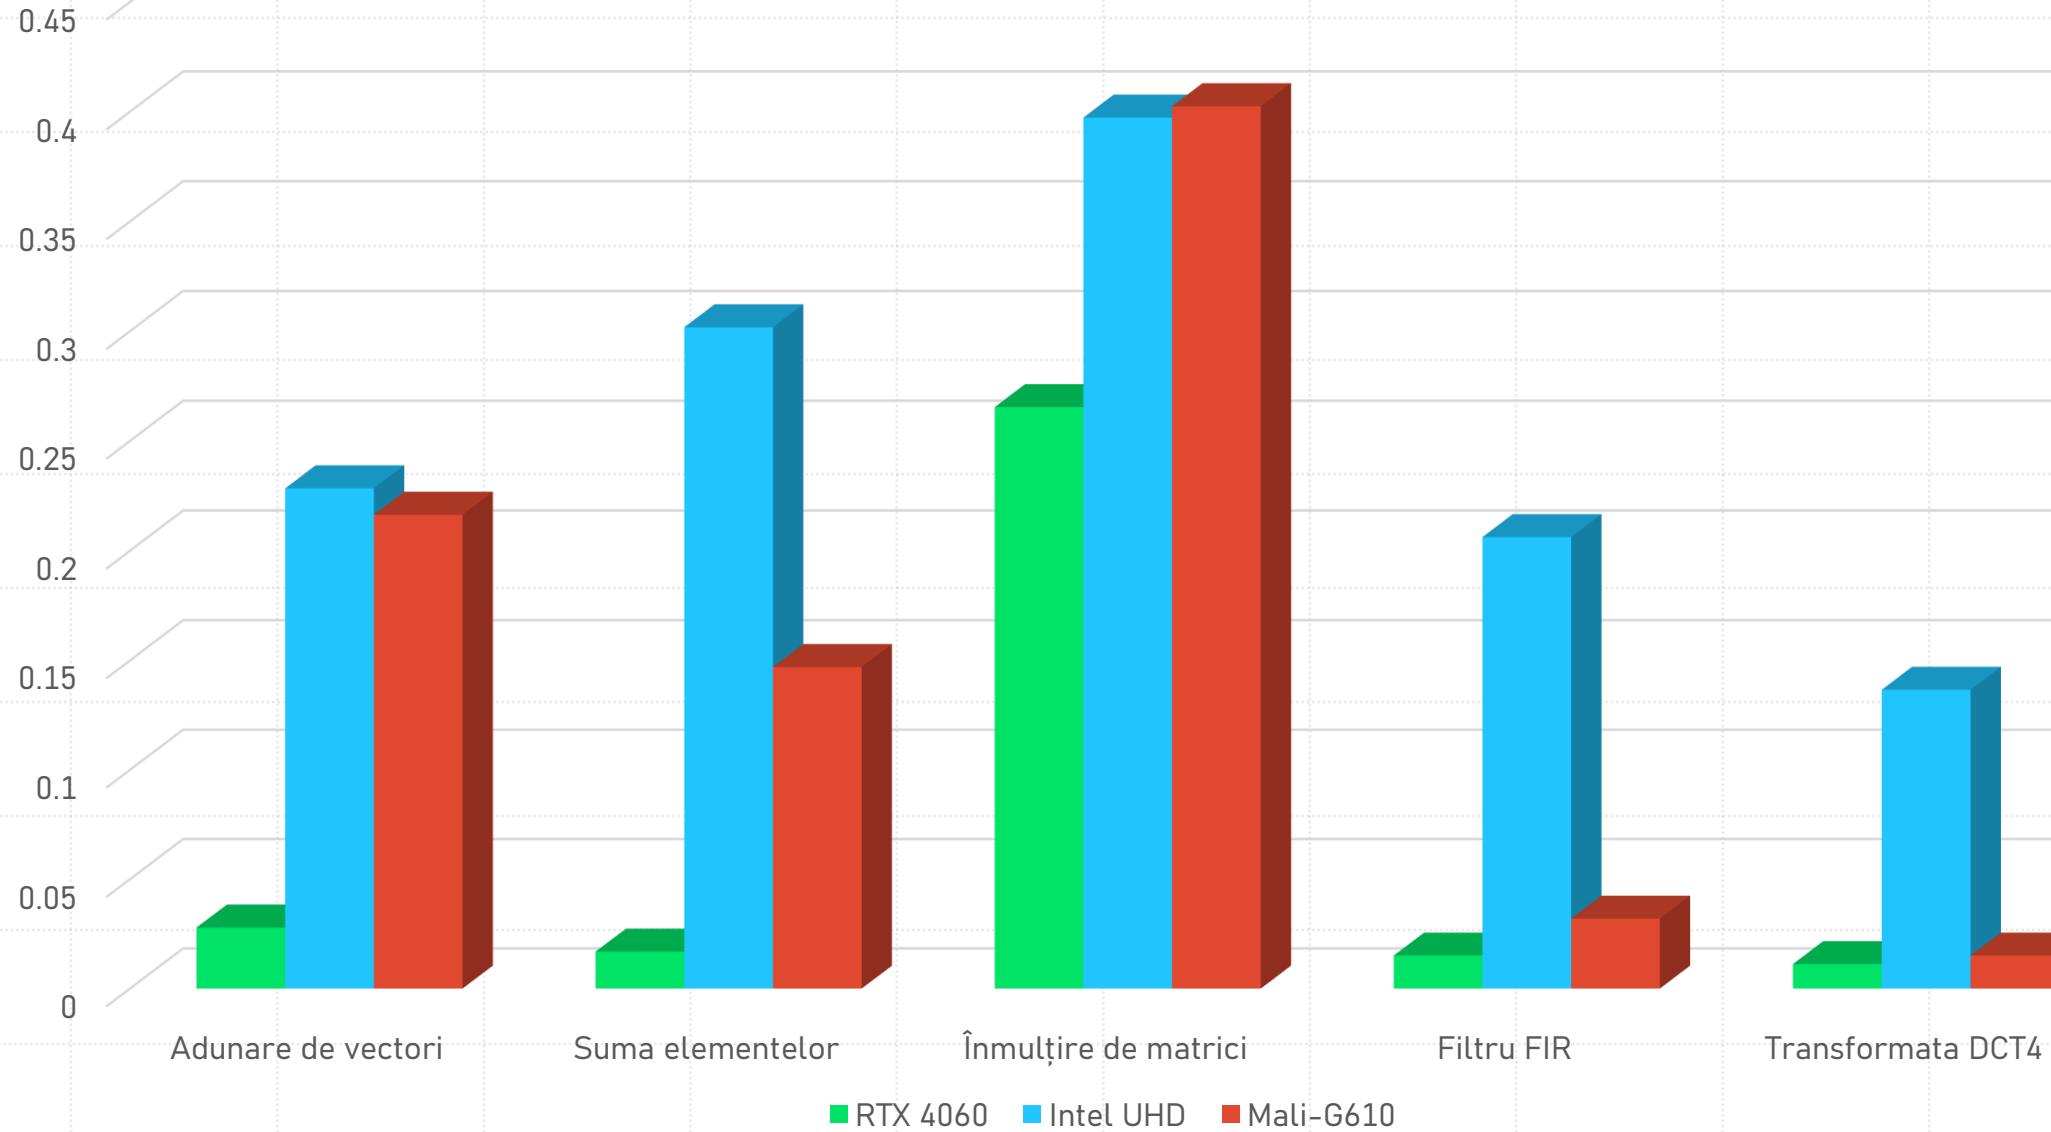
# Algoritmii folosiți pentru testarea performanței

- Adunare de vectori (2 vectori cu 4096 de elemente)
- Suma elementelor dintr-un vector (4096 de elemente)
- Înmulțire de matrici (Matrice 1: 256x512, Matrice 2: 512x128)
- Filtru FIR (Bloc de intrare cu 384 de valori și 256 coeficienți) \*
- Transformata DCT4 (Vector de dimensiune 2048) \*

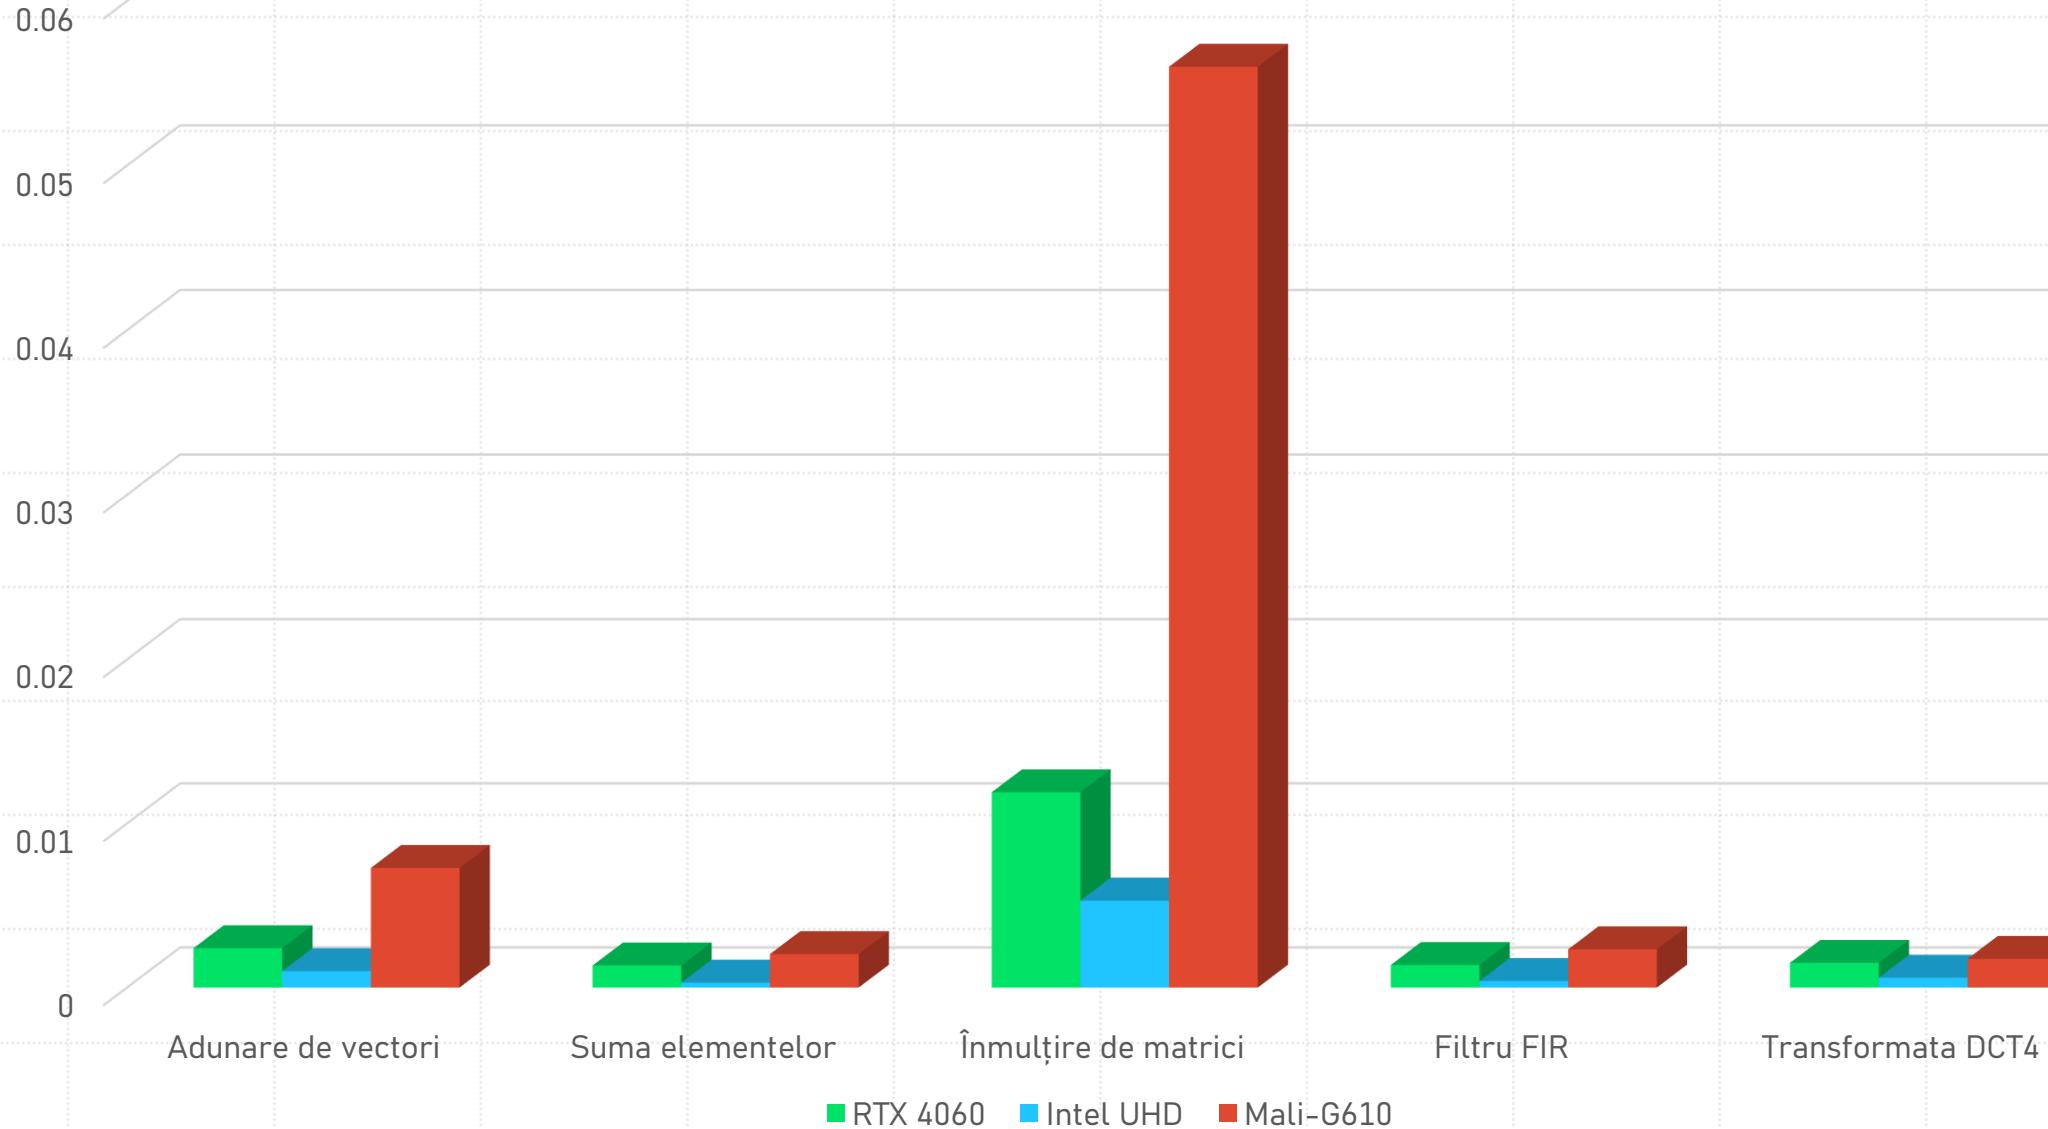
Toate elementele acestor vectori și matrici sunt float32.

\* Acești algoritmi au fost luați din articolul: “Embench IOT 2.0 and DSP 1.0: Modern Embedded Computing Benchmarks”

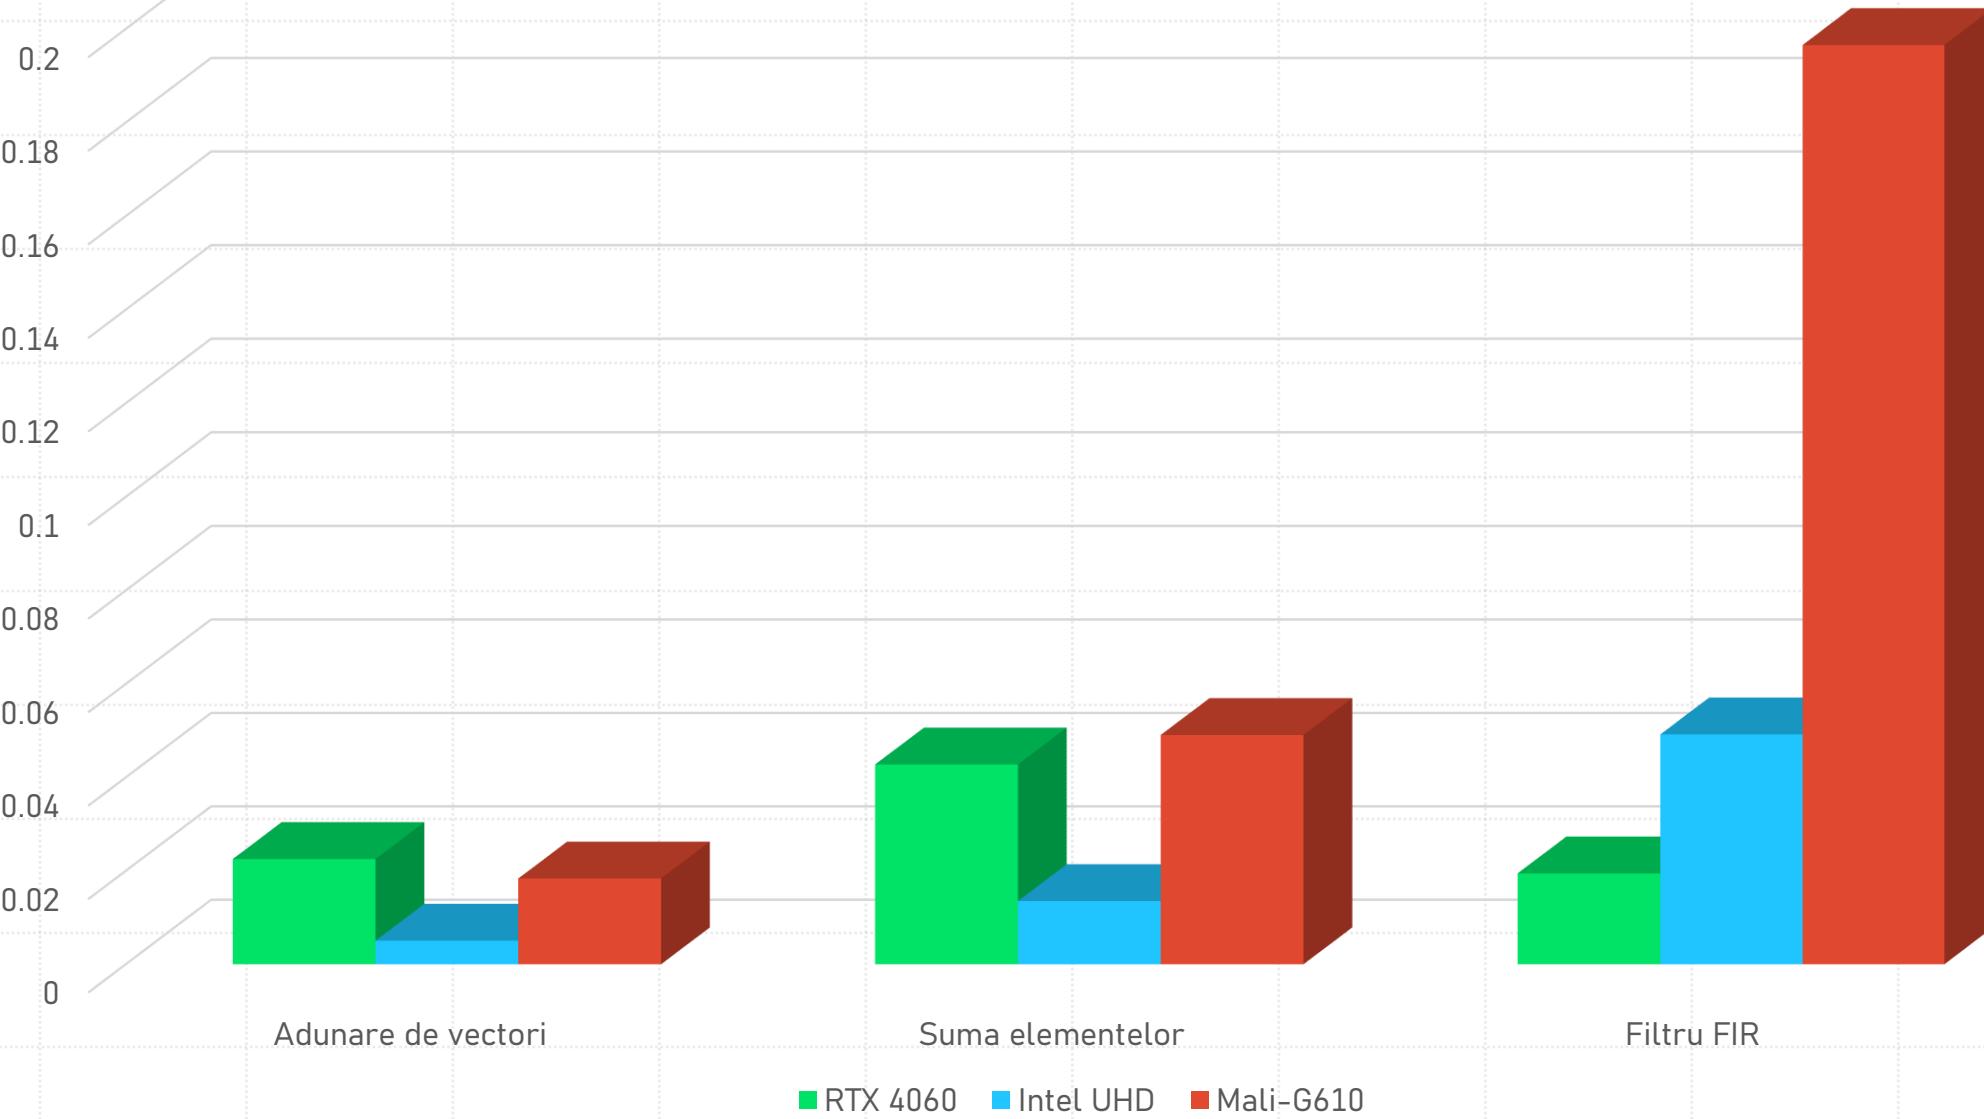
## Timpul de scriere în buffer (în milisecunde)



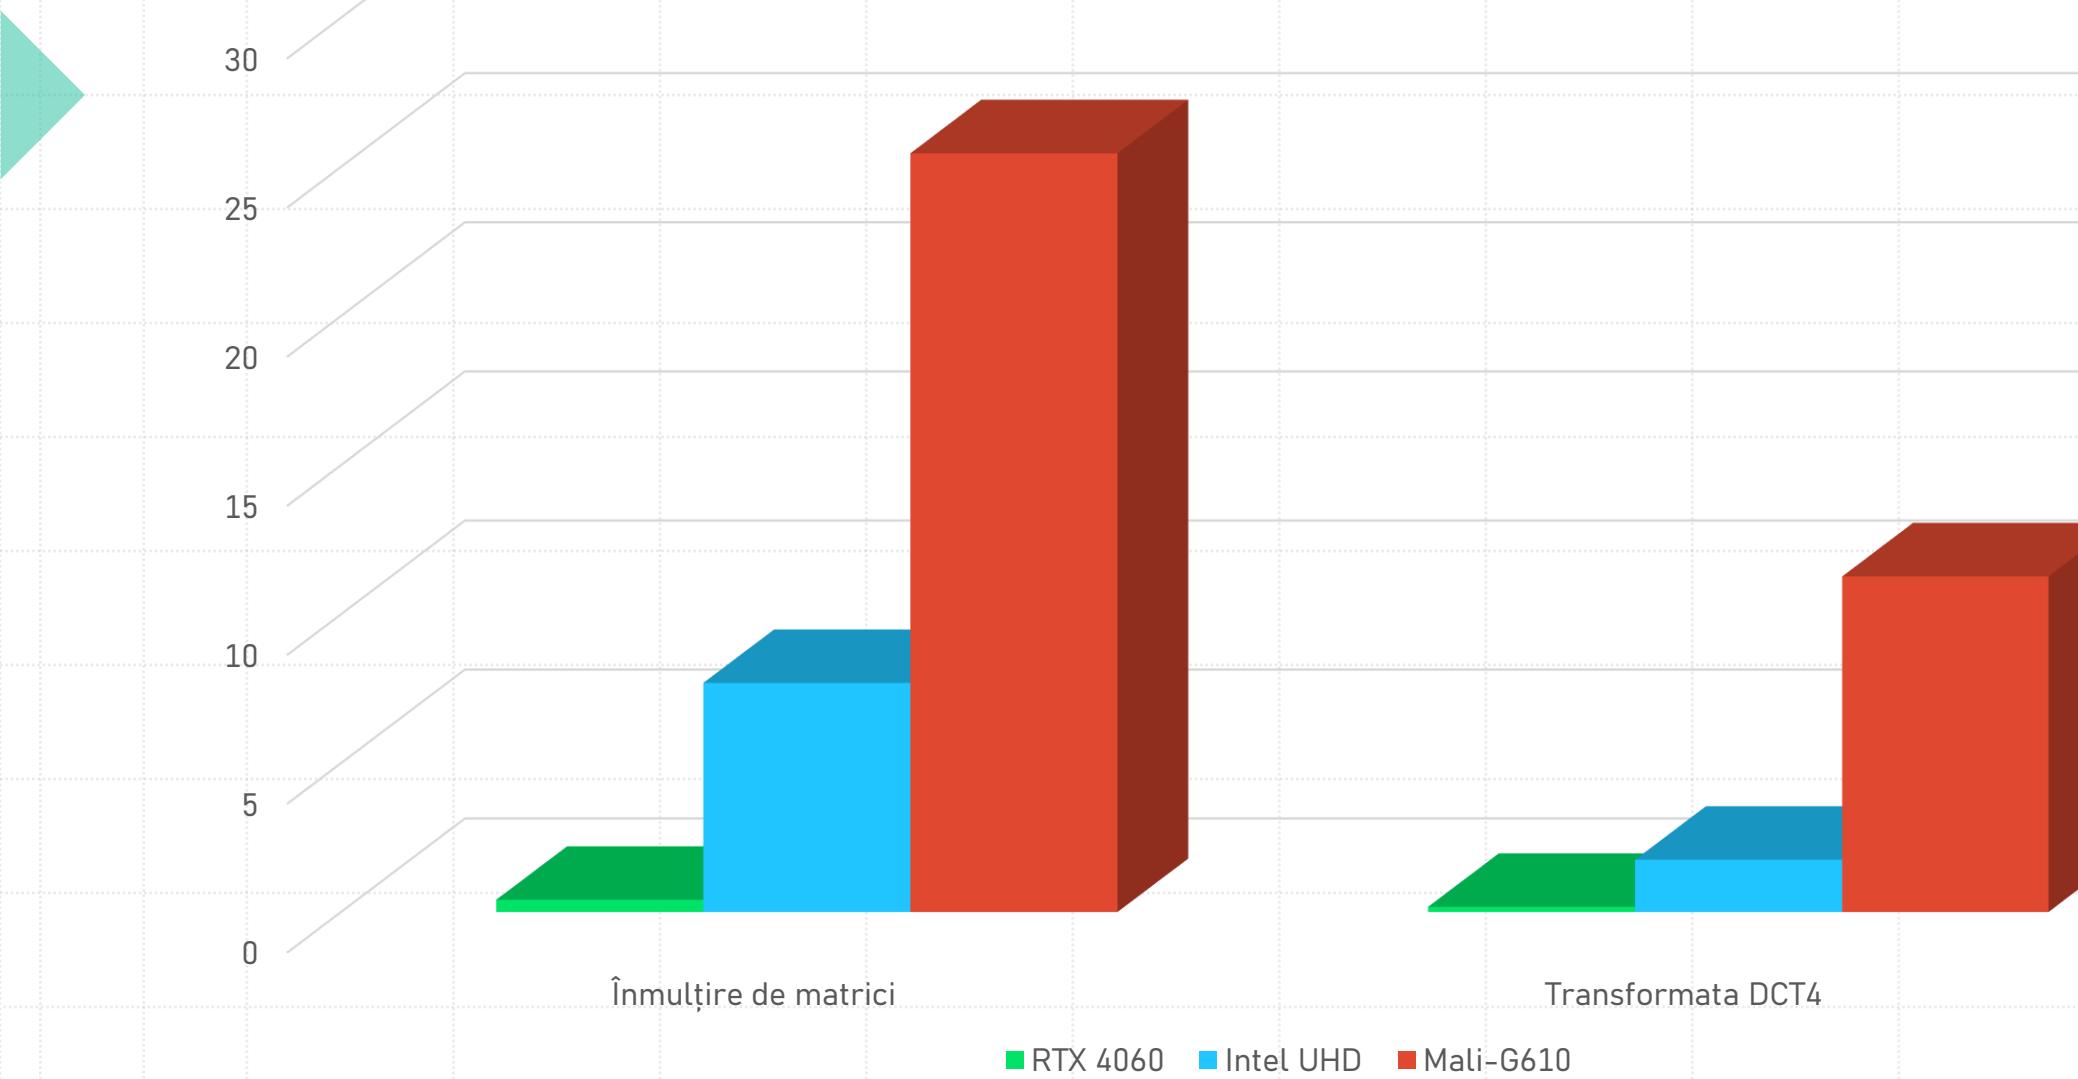
## Timpul de citire din buffer (în milisecunde)



## Timp de execuție al funcției kernel (în milisecunde)



## Timp de execuție al funcției kernel (în milisecunde) - continuare



# Concluzii și direcții viitoare

## Concluzii

- Framework-ul simplifică dezvoltarea și testarea aplicațiilor OpenCL
- Este modular, portabil și compatibil cu mai multe tipuri de dispozitive
- Suportă execuția de funcții kernel înlănțuite

## Direcții viitoare de dezvoltare

- Execuția distribuită automată pe mai multe dispozitive
- Mecanism de fallback
- Interfață grafică mai amplă



# Întrebări?